

Survey of Machine Learning Feature Selection Methods: Boruta, SVD, PCA, LASSO, RFE, ...

Earl F Glynn
Kansas City R Users Group
2018-09-08

<https://github.com/EarlGlynn/kc-r-users-feature-selection>

Continuation from last year ...



Using R's Caret Package for Machine Learning

Earl F Glynn

Kansas City R Users Group

9 September 2017

<https://github.com/EarlGlynn/kc-r-users-caret-2017>

From last year

Caret Machine Learning Examples

Group	Algorithms	Caret Model	FILE
Linear Methods	Linear Discriminant Analysis	lda	lda
	Linear Discriminant Analysis with YeoJohnson preprocessing	lda w/YeoJohnson	lda-YeoJohnson
	LASSO, Ridge, and Elastic Net	glmnet	glmnet
	LASSO, Ridge, and Elastic Net with Synthetic Minority Over-Sampling Technique (SMOTE)	glmnet w/SMOTE	glmnet-SMOTE
Non-Linear Methods	Neural Network	nnet	nnet
	Support Vector Machine with Radial Basis Function Kernel	svmRadial	svmRadial
	Naïve Bayes	nb	nb
	Naïve Bayes with Independent Component Analysis (ICA)	nb w/ICA	nb-ica-transform
	k Nearest Neighbors	knn	knn
Trees and Rules	J48	J48	J48
	Classification and Regression Trees	rpart	rpart
Ensembles of Trees	C5.0	C5.0	C50
	Random Forest	rf	rf
	Random Forest with SMOTE	rf w/SMOTE	rf-SMOTE

<https://github.com/EarlGlynn/kc-r-users-caret-2017>

Files: **Forensic-Glass-caret-FILE.Rmd** and **Forensic-Glass-caret-FILE.html**

Survey of ...

FEATURE SELECTION STRATEGIES

1. Remove highly correlated variables.
2. Run OLS and select significant features.
3. Forward selection and backwards selection. ↘
4. Random Forest feature importance. or recursive
5. Lasso.

BY CHRIS ALBON

6. Boruta "All Relevant" Variables
7. SVD (Singular Value Decomposition)
8. PCA (Principal Component Analysis)

Small dataset used in examples: Forensic Glass Dataset




<https://sha.org/bottle/links.htm>

fgl dataset in MASS package

- RI = refractive index
- Percentages by weight of oxides: Na, Mg, Al, Si, K, Ca, Ba, Fe
- Predict *type* from measured values above
 - window float glass (WinF)
 - window non-float glass (WinNF)
 - vehicle window glass (Veh)
 - containers (Con)
 - tableware (Tabl)
 - vehicle headlamps (Head)

<https://toolbox.google.com/datasetsearch>



Also available through [UCI Repository](#).

Discussed in book [Data Mining and Business Analytics with R](#).

Forensic Glass Dataset

Forensic Glass Data

```
library(MASS)      # fgl data
```

```
rawData <- fgl  
dim(rawData)
```

```
[1] 214  10
```

```
rawData  
kable("html", caption="Forensic Glass Data")  
kable_styling(bootstrap_options=c("striped", "bordered", "condensed"),  
              position="left", font_size=12,  
              full_width=FALSE)  
scroll_box(height="200px")
```

Forensic Glass Data

RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	type
3.01	13.64	4.49	1.10	71.78	0.06	8.75	0.00	0.00	WinF
-0.39	13.89	3.60	1.36	72.73	0.48	7.83	0.00	0.00	WinF
-1.82	13.53	3.55	1.54	72.99	0.39	7.78	0.00	0.00	WinF
-0.34	13.21	3.69	1.29	72.61	0.57	8.22	0.00	0.00	WinF
0.58	13.37	3.63	1.34	72.88	0.55	8.07	0.00	0.00	WinF

```
table(rawData$type)
```

```
WinF WinNF  Veh  Con  Tabl  Head  
  70   76   17   13    9   29
```

Class imbalance

Files: **Forensic-Glass-Boruta.Rmd** and **Forensic-Glass-Boruta.html**

Feature Selection

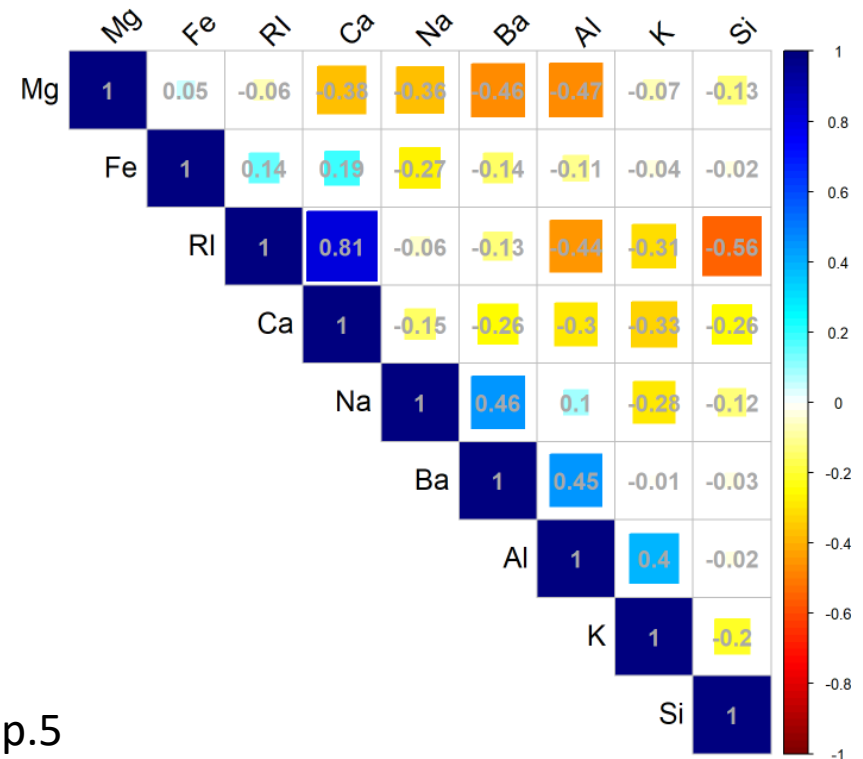
1. Remove Highly Correlated Variables

trainSet Correlation Matrix

Use only training set data or we'll have a **data leak**.

```
colorScale <- colorRampPalette(c("#7F0000", "red", "#FF7F00", "yellow", "white",  
                                "cyan", "#007FFF", "blue", "#00007F"))(100)  
  
corMatrix <- cor(trainSet %>% dplyr::select(-type)) # Create correlation matrix  
  
corrplot(corMatrix, type="upper", method="square", order="AOE",  
          tl.col="black", tl.srt=45, tl.cex=1.5,  
          addCoef.col="darkgrey", number.cex=1.25,  
          col=colorScale)  
mtext("Correlation Matrix (angular order of eigenvectors)", line=3)
```

Correlation Matrix (angular order of eigenvectors)



order="AOE" (angular order of eigenvalues)
<http://www.datavis.ca/papers/corrgram.pdf>, p.5

Feature Selection

Remove Highly Correlated Variables

I often use a cutoff of 0.9 or above, but here 0.8 is used to trigger a removal.

Removing highly correlated pairs with caret

Some machine learning algorithms are impeded by highly correlated predictors. Caret's `findCorrelation` procedure can be used to remove one of the highly-correlated pairs.

```
HIGH_CORRELATION_CUTOFF <- 0.80

corHigh <- findCorrelation(corMatrix, HIGH_CORRELATION_CUTOFF)
if (length(corHigh) > 0)
{
  cat("Removing highly-correlated variable(s): ", names(trainSet)[corHigh])
  trainSet <- trainSet[, -corHigh]
  testSet  <- testSet[, -corHigh]
}
```

```
Removing highly-correlated variable(s):  Ca
```

Note: Remove constant variables, too.

Files: **Forensic-Glass-Correlation.Rmd** and **Forensic-Glass-Correlation.html**

Feature Selection

2. Run OLS and select significant features

OLS = Ordinary Least Square, but approach also applies to
NLS = Non-linear Least Squares, e.g.,
Levenberg-Marquardt non-linear curve fitting

```
library(minpack.lm) # nls.lm
```

```
summary(nlsFit)
```

Parameters:

	Estimate	Std. Error	t value	Pr(> t)
A	29.518818	23.888092	1.236	0.263
alpha	-4.615019	3.485175	-1.324	0.234
B	86.974831	2.210565	39.345	1.80e-08 ***
beta	-0.092361	0.005677	-16.269	3.43e-06 ***

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 2.002 on 6 degrees of freedom

Number of iterations to termination: 9

Reason for termination: Relative error in the sum of squares is at most 'ftol'.

Feature Selection

3. Caret's Recursive Feature Extraction (RFE)

Random Forests

```
rfeController <- rfeControl(functions = rfFuncs, # random forests
  seeds = seeds,
  method = "repeatedcv",
  number = nCV,
  repeats = nRepeats,
  verbose = FALSE)
```

Setup parallel processing

```
rCluster <- makePSOCKcluster(6) # use 6 cores
registerDoParallel(rCluster)

rfProfile <- rfe(x, y, sizes=subsets, rfeControl=rfeController)

stopCluster(rCluster)
```

rfProfile

Recursive feature selection

Outer resampling method: Cross-Validated (3 fold, repeated 50 times)

Resampling performance over subset size:

Variables	Accuracy	Kappa	AccuracySD	KappaSD	Selected
1	0.3860	0.1740	0.05975	0.07562	
2	0.6170	0.4725	0.07747	0.10338	
3	0.6670	0.5383	0.05204	0.07326	
4	0.7048	0.5954	0.06086	0.08318	
5	0.7300	0.6306	0.06245	0.08546	
6	0.7507	0.6583	0.05779	0.07939	
7	0.7647	0.6779	0.05450	0.07429	*
8	0.7544	0.6622	0.05308	0.07298	
9	0.7581	0.6669	0.05218	0.07218	

The top 5 variables (out of 7):

Mg, Al, Ba, RI, Na

Output shows the best subset size was 7 predictors.

predictors(rfProfile)

[1] "Mg" "Al" "Ba" "RI" "Na" "Ca" "K"

-Si -Fe

Naive Bayes

```
rfeController <- rfeControl(functions = nbFuncs, # naive bayes
  seeds = seeds,
  method = "repeatedcv",
  number = nCV,
  repeats = nRepeats,
  verbose = FALSE)
```

Setup parallel processing

```
rCluster <- makePSOCKcluster(6) # use 6 cores
registerDoParallel(rCluster)

nbProfile <- rfe(x, y, sizes=subsets, rfeControl=rfeController)

stopCluster(rCluster)
```

nbProfile

Recursive feature selection

Outer resampling method: Cross-Validated (3 fold, repeated 50 times)

Resampling performance over subset size:

Variables	Accuracy	Kappa	AccuracySD	KappaSD	Selected
1	0.5020	0.2889	0.08511	0.12103	
2	0.6093	0.4544	0.05870	0.08073	
3	0.6274	0.4888	0.06283	0.08532	
4	0.6311	0.4994	0.06330	0.08583	
5	0.6358	0.5087	0.06242	0.08254	*
6	0.6252	0.4944	0.06187	0.08170	
7	0.6234	0.4931	0.06358	0.08370	
8	0.6046	0.4663	0.06458	0.08685	
9	0.5846	0.4358	0.06598	0.08868	

The top 5 variables (out of 5):

Al, Mg, Na, Ca, K

Output shows the best subset size was 5 predictors.

predictors(nbProfile)

[1] "Al" "Mg" "Na" "Ca" "K"

-RI -Si -Ba -Fe


Files: **Forensic-Glass-caret-RFE.Rmd** and **Forensic-Glass-caret-RFE.html**

Feature Selection

Caret's Recursive Feature Extraction (RFE)


```
plot(rfProfile, type=c("g", "o"), metric="Kappa",  
     main="rf: Kappa by Number of Variables")
```

rf: Kappa by Number of Variables



```
plot(nbProfile, type=c("g", "o"), metric="Kappa",  
     main="nb: Kappa by Number of Variables")
```

nb: Kappa by Number of Variables



Kappa is generally a better metric than *Accuracy* with imbalanced classification.

Files: **Forensic-Glass-RFE.Rmd** and **Forensic-Glass-RFE.html**

Feature Selection

4. Feature Importance

FEATURE IMPORTANCE

Decision trees make splits that maximize the decrease in impurity.

By calculating the mean decrease in impurity for each feature across all trees we can know that feature's importance.




ChrisAlbon

Feature Selection


4. Feature Importance

Use feature importance from one or more machine learning models to select features.
Study consistency and differences among models. Single Predictors here.

Tree Method
Random Forest Classification
Variable Importance



Ensemble of Trees Method
Classification and Regression Trees
Variable Importance




2017 Files: **Forensic-Glass-caret-rf.html** and **Forensic-Glass-caret-rpart.html**

Feature Selection


Feature Importance

Use feature importance from one or more machine learning models to select features.
Study consistency and differences among models. Separate Predictors, Very Similar.

Linear Method Linear Discriminant Analysis



Non-Linear Method Naïve Bayes Classifier



2017 Files: **Forensic-Glass-caret-lda.html** and **Forensic-Glass-caret-nb.html**

Feature Selection

5. glmnet

Ridge, Lasso and Elastic-Net Regularized Generalized Linear Models
($\alpha=0$ for Ridge, $\alpha=1$ for Lasso)

- Math is a bit complicated:

<https://quantmacro.wordpress.com/2016/04/26/fitting-elastic-net-model-in-r/>
https://web.stanford.edu/~hastie/glmnet/glmnet_alpha.html

- Idea is to force some coefficients to zero to exclude from model.

REGULARIZATION

A variety of techniques used to improve generalizability by changing the learning algorithm to improve performance on the test error, regardless of the effect on the training error.

EXAMPLES

- Weight decay
- Drop-out
- Ridge regression
- Lasso regression

Chris Albon

Feature Selection

glmnet

glmnet (lasso and elastic-net regularization)

```
tuneGrid <- expand.grid(alpha = seq(0.25, 0.75, by=0.05), # alpha 1 for Lasso, 0 for Ridge
                      lambda = c(0.05, 0.005, 0.0005)) # strength of penalty on coefficients

set.seed(29)
CVfolds <- 5 # 5-fold cross validation (not enough data for 10 fold here)
CVrepeats <- 10 # repeat 10 times

# Used createMultiFolds to study
indexFolds <- createMultiFolds(trainSet$type, CVfolds, CVrepeats) # for repeated CV

trainControlParms <- trainControl(method = "repeatedcv", # repeated cross validation
                                number = CVfolds,
                                repeats = CVrepeats,
                                index = indexFolds,
                                classProbs = TRUE, # Estimate class probabilities
                                summaryFunction = defaultSummary)

fit <- train(type ~ ., data=trainSet,
            preProcess = c("center", "scale"),
            method = "glmnet",
            metric = "Kappa",
            tuneGrid = tuneGrid,
            trControl = trainControlParms)
```

Forensic-Glass-caret-glmnet.Rmd and Forensic-Glass-caret-glmnet.html

Feature Selection

glmnet

Each class has a separate set of coefficients.

Fit coefficients for best λ :

```
coef(fit$finalModel, s = fit$bestTune$lambda)
```

```
$WinF
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)  1.207375322
RI           0.250065503
Na          -0.224390621
Mg           1.630701058
Al          -1.402125023
Si           0.004345775
K            .
Ca            .
Ba           0.277709942
Fe            .

$WinNF
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)  2.0909299
RI           .
Na          -0.7154609
Mg           0.1743042
Al           .
Si          -0.9109666
K           0.1107151
Ca          -0.1779174
Ba           .
Fe           0.2913854

$Veh
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)  0.16121177
RI          -1.02662928
Na           .
Mg           0.56240805
Al          -1.36706675
Si          -1.29460588
K           -0.48120821
Ca           0.05261939
Ba           .
Fe           0.14586035
```

```
$Con
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) -1.5639444
RI           .
Na          -0.6234961
Mg          -1.3894046
Al           1.4634459
Si           0.1103138
K            1.0771763
Ca           0.5982293
Ba          -0.5683051
Fe           0.4363922

$Tabl
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) -1.1941202
RI           .
Na           1.4674086
Mg          -0.1975055
Al           .
Si           0.4315638
K           -1.7011573
Ca           0.1962890
Ba          -0.6153015
Fe          -0.6018880

$Head
10 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) -0.70145244
RI           0.86756205
Na          1.23281956
Mg          -0.78050324
Al           1.29083825
Si           0.02129146
K            .
Ca          -0.82978126
Ba           1.29411893
Fe          -0.65922012
```

glmnet creates a separate model for each class.

Some coefficients are driven to zero by the method, which effectively excludes predictor from model.

glmnet can be used as a prediction model, or as a tool to select relevant variables for other models.

Forensic-Glass-caret-glmnet.Rmd and Forensic-Glass-caret-glmnet.html

Feature Selection

6. Boruta “All Relevant” Variables

Define train and test datasets

```
set.seed(71)

trainSetIndices <- createDataPartition(rawData$type, p=0.50, list=FALSE)

trainSet <- rawData[ trainSetIndices, ]
testSet  <- rawData[-trainSetIndices, ]
```

```
dim(trainSet)
```

```
## [1] 109 10
```

Boruta ‘All Relevant’ Variables

Setup parallel processing

```
rCluster <- makePSOCKcluster(6) # use 6 cores
registerDoParallel(rCluster)
```

Will vary
by platform

```
set.seed(73) # reproducible random numbers
BorutaModel <- Boruta(type ~ ., data=trainSet,
                      getImp = getImpForns,
                      maxRuns = 1000, num.threads=6)
stopCluster(rCluster)
```

No data
pre-processing
needed!



Boruta was a demon/god in Slavic mythology, who lived in trees.

<https://www.pinterest.com/pin/497577458807361039/>

Feature Selection

Boruta “All Relevant” Variables

```
print(BorutaModel)
```

```
Boruta performed 48 iterations in 1.187219 secs.  
8 attributes confirmed important: Al, Ba, Ca, K, Mg and 3 more;  
1 attributes confirmed unimportant: Fe;
```

Basic Idea of Boruta Algorithm

Perform shuffling of predictors' values and join them with the original predictors and then build random forest on the merged dataset. Then make comparison of original variables with the randomised variables to measure variable importance. Only variables having higher importance than that of the randomised variables are considered important.

<https://www.listendata.com/2017/05/feature-selection-boruta-package.html>

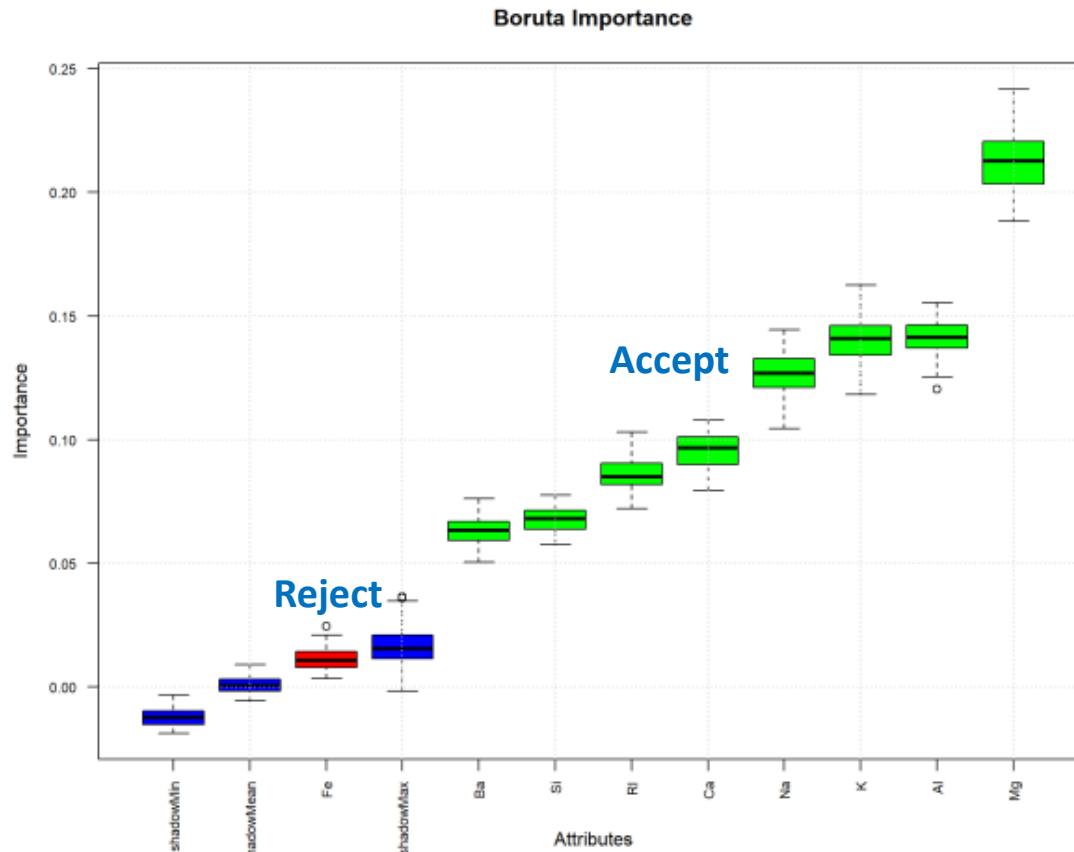
<https://www.datacamp.com/community/tutorials/feature-selection-R-boruta>

Files: **Forensic-Glass-Boruta.Rmd** and **Forensic-Glass-Boruta.html**

Feature Selection

Boruta “All Relevant” Variables

```
plot(BorutaModel, las=2, cex.axis=0.75, main="Boruta Importance")  
grid()
```

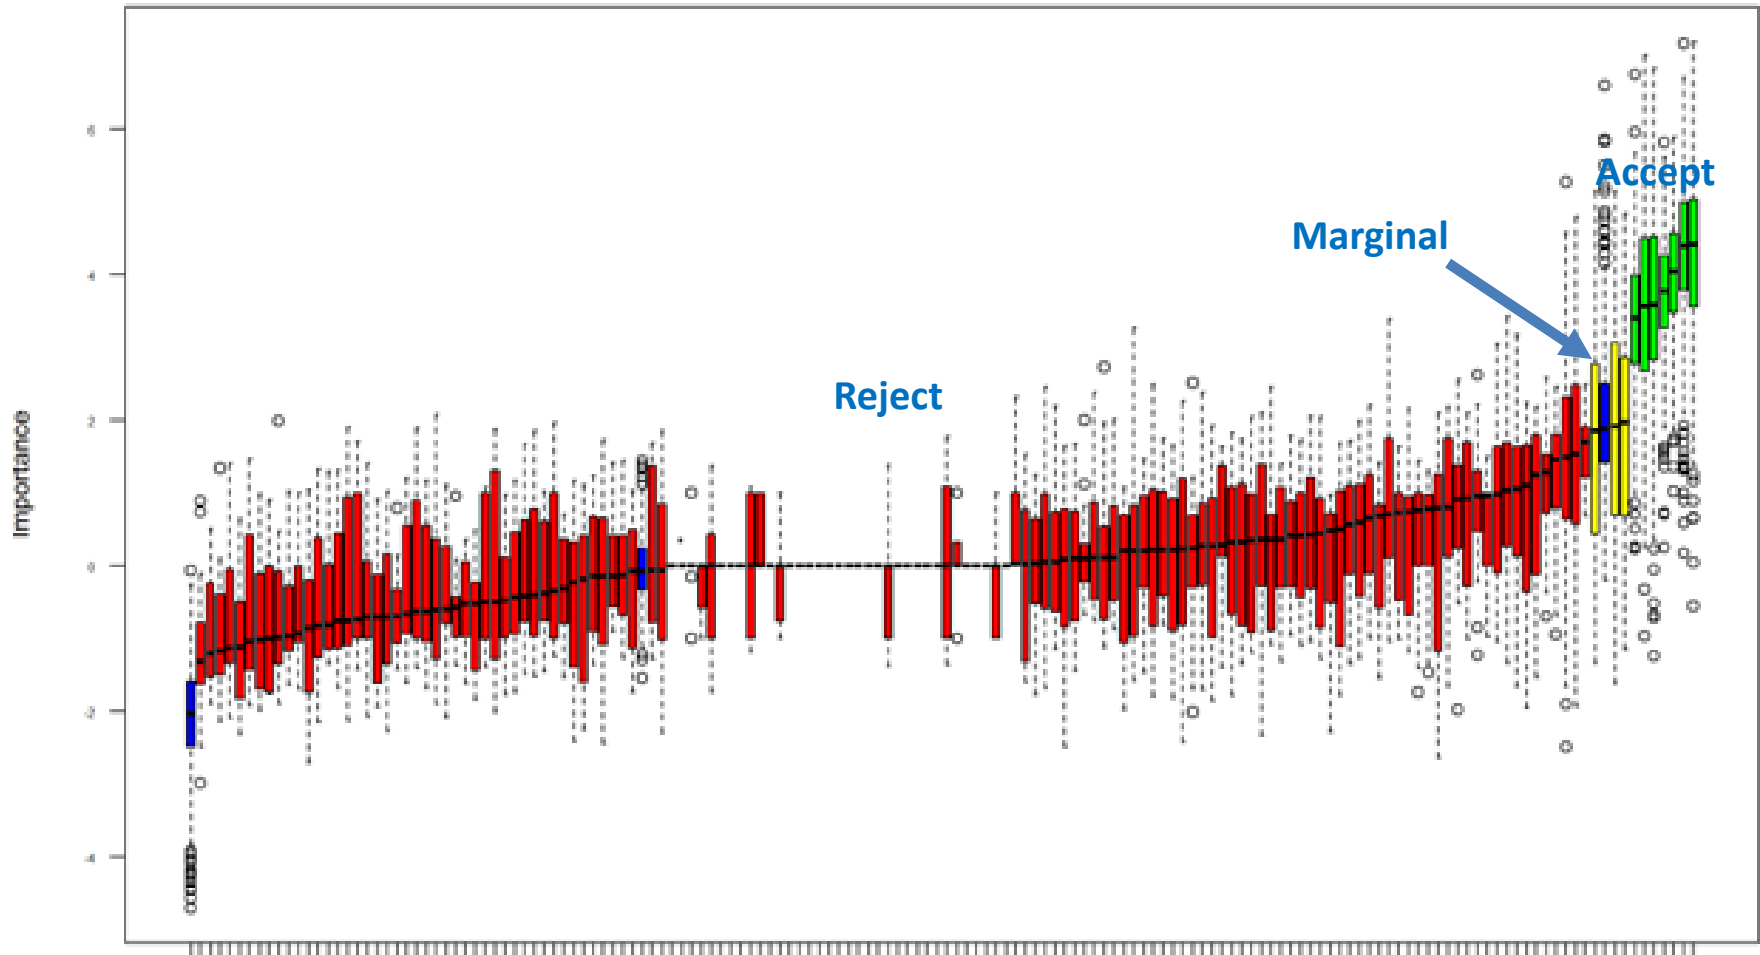


Files: **Forensic-Glass-Boruta.Rmd** and **Forensic-Glass-Boruta.html**

Feature Selection

Boruta “All Relevant” Variables

Sometimes Boruta approach can be very helpful with selections



Feature Selection

Boruta “All Relevant” Variables

```
stats <- attStats(BorutaModel)
statsOrdered <- stats[order(stats$medianImp, decreasing=TRUE), ]
```

```
BorutaFeatures <- statsOrdered %>% rownames_to_column(var="Feature")

BorutaFeatures %>%
  kable("html", caption="Boruta Features") %>%
  kable_styling(bootstrap_options=c("striped", "bordered", "condensed"),
    position="left", font_size=12,
    full_width=FALSE)
```

Boruta Features

Feature	meanImp	medianImp	minImp	maxImp	normHits	decision
Mg	0.2127641	0.2125817	0.1883694	0.2417506	1.0000000	Confirmed
Al	0.1406346	0.1412464	0.1204132	0.1553545	1.0000000	Confirmed
K	0.1400205	0.1408165	0.1181545	0.1625252	1.0000000	Confirmed
Na	0.1260519	0.1267271	0.1045717	0.1443913	1.0000000	Confirmed
Ca	0.0953332	0.0965700	0.0795888	0.1080054	1.0000000	Confirmed
Rl	0.0860814	0.0849388	0.0722158	0.1030754	1.0000000	Confirmed
Si	0.0676871	0.0681237	0.0576441	0.0776154	1.0000000	Confirmed
Ba	0.0630710	0.0633130	0.0504127	0.0762487	1.0000000	Confirmed
Fe	0.0113589	0.0108034	0.0033247	0.0246047	0.2708333	Rejected

Feature Selection

Boruta “All Relevant” Variables

```
plotImpHistory(BorutaModel, main="Importance History")  
grid()
```



Feature Selection

7. Singular Value Decomposition

- Dataset is a matrix of 214 rows of 9 predictors.
- Original matrix can be written as a product of three matrices with svd.
- Diagonal of middle matrix contains the singular values (eigenvalues if centered data).
- Right-most matrix contains the right singular vectors, which can give information about feature importance.

$$\text{Train Matrix} = \mathbf{U} \times \mathbf{\Sigma} \times \mathbf{V}^T$$

In R:

```
svd1 <- svd(scale(trainData))  
svd1$u %*% diag(svd1$d) %*% t(svd1$v)
```

$$\text{Train} = \mathbf{U} \times \mathbf{\Sigma} \times \mathbf{V}^T$$

Diagram illustrating the dimensions of the matrices in the SVD equation:

- \mathbf{U} is an $n \times k$ Orthogonal matrix.
- $\mathbf{\Sigma}$ is a $k \times k$ Diagonal matrix.
- \mathbf{V}^T is a $k \times k$ Orthogonal matrix.

\mathbf{U} = left singular vectors = 214-by-9 unitary matrix. **svd1\$u** in R.

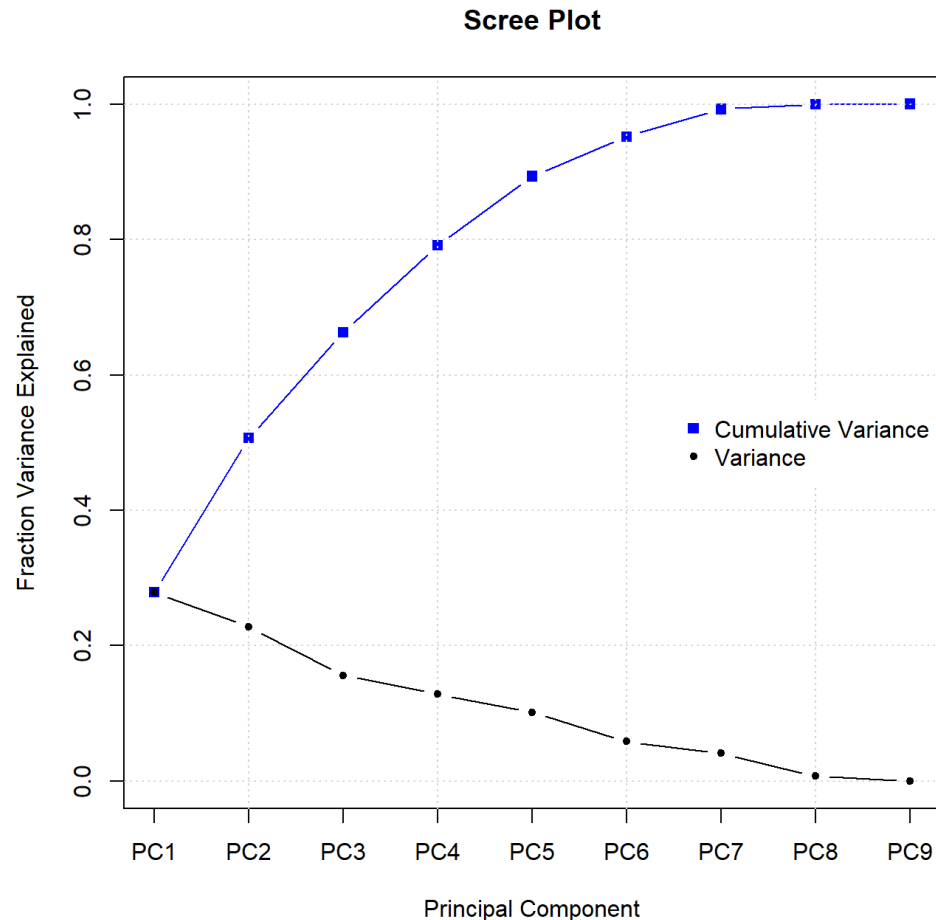
$\mathbf{\Sigma}$ = singular values matrix = 9-by-9 diagonal matrix. **diag(svd1\$d)** in R.

The diagonal terms are the singular values, usually listed in decreasing order.

\mathbf{V}^T = transpose of \mathbf{V} , where **\mathbf{V} = right singular vectors** = 9x9 unitary matrix. **svd1\$v** in R.

Feature Selection

Singular Value Decomposition



Singular values are usually in decreasing order. 1st one largest.

Square of n^{th} singular value proportional to variance associated with n^{th} singular vector.

Often few terms explain large % of total variance.

Cumulative variance explained:

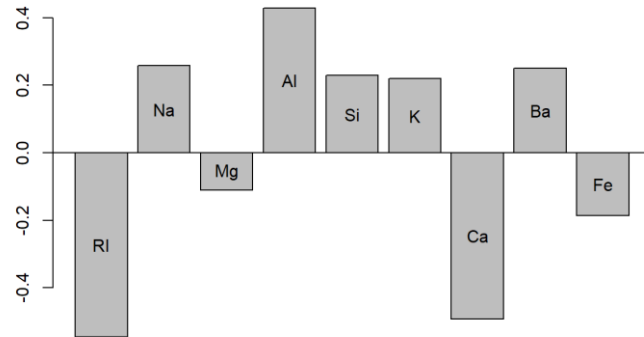
```
[1] 0.2790 0.5068 0.6629 0.7915  
    0.8931 0.9517 0.9927 0.9998  
    1.0000
```

Files: **Forensic-Glass-SVD.Rmd** and **Forensic-Glass-SVD.html**


Feature Selection

Singular Value Decomposition


28% of variance Right singular vector 1




23% of variance Right singular vector 2



16% of variance Right singular vector 3



13% of variance Right singular vector 4




Pick variables with large contributions (+ or -): 1: RI, Ca, Al; 2: Mg, Ba; 3: K, Si

Feature Selection

8. Principal Component Analysis


PCA often computed from SVD



Feature Selection

SVD / PCA

Samsung



Right singular vector can help with interpretation of a principal component.
Here, 9 of 500+ PCs explained over 70% of total variance.

Could use 9 PCs in machine learning that explain over 70% of variance
instead of 500+ predictors.

Feature Selection

Principal Component Analysis

Setup

NOTE: place `tidyverse` after `MASS` below to avoid `dplyr::select` function clashes with `MASS::select`.

```
library(MASS)           # fgl data
library(tidyverse)      # place after MASS to avoid select conflict
library(caret)          # preProcess, predict
library(rgl)            # par3d, plot3d, movie3d, rglwidget
library(RColorBrewer)   # brewer.pal
```

Principal Component Analysis

Let's compute the values for the first 4 principal components using caret's pca pre-processing.

"pca" requires "center" and "scale".

These four PCs account for nearly 80% of variance.

```
nPcAcomponents <- 4
transformSetup <- preProcess(rawData, method=c("center", "scale", "pca"), pcaComp=nPcAcomponents)
pcaScores <- predict(transformSetup, rawData)
```

```
pcaScores %>% head(10)
```

	PC1	PC2	PC3	PC4
1	-1.148446843	-0.5282491	0.3712253	-1.72485681
2	0.572794160	-0.7580105	0.5554059	-0.75845396
3	0.937960515	-0.9276609	0.5536094	-0.20577184
4	0.141750924	-0.9594279	0.1168507	-0.41475157
5	0.350271021	-1.0886966	0.4839440	-0.06894065

Files: **Forensic-Glass-PCA.Rmd** and **Forensic-Glass-PCA.html**

Feature Selection

Principal Component Analysis

- *caret's preProcess* gives the same PCAscores as computed with SVD.
- Each PC is a weighted linear combination of all variables.
- PCs are orthogonal.
- **PCs can be used as variables in other machine learning algorithms.**
- Machine learning algorithm using PCs as predictors are limited by the amount of variance explained by the original variables in the given number of PCs.

Interactive Exploratory Analysis

Principal Component Analysis

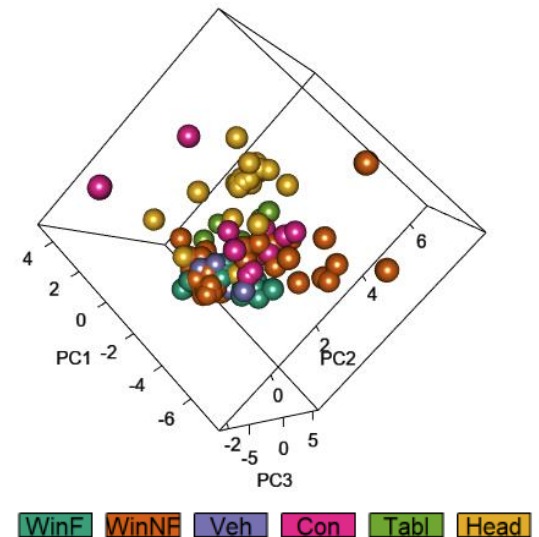
Use **plot3d** in **rgl** package to create interactive 3D scatterplot of any three PCs

The first 3 PCs account for about 66% of variance in data.

Other sets of 3 PCs could be displayed alternatively in 3D space, such as PC2, PC3, PC4.

```
typeColors <- brewer.pal(length(levels(fgl$type)), "Dark2")
```

```
par3d("windowRect"=c(50,50,800,800))  
plot3d(x=pcaScores$PC1, y=pcaScores$PC2, z=pcaScores$PC3,  
       col=typeColors[typeColorIndex],  
       xlab="PC1", ylab="PC2", zlab="PC3", type="s", size= 3)  
rglwidget(elementId="FGL1")
```



Chrome browser works best to display above figure.

Drag mouse over figure to rotate. Use mouse wheel to zoom in and out.

Files: **Forensic-Glass-PCA.Rmd** and **Forensic-Glass-PCA.html**

Exploratory Analysis

Principal Component Analysis

magick from **ImageMagick** needed to create animated GIF of 3D PCA scatterplot

Animated GIF

Create the animated GIF using **magick** from **ImageMagick** – this takes some time. Display below using HTML.

150 PNG images will be computed for 15 sec duration * 10 frames/second.

```
movie3d(spin3d(), duration = 15, dir = getwd(),  
        movie="ForensicGlass-PCA",  
        verbose=FALSE, convert="magick -delay 1x%d %s*.png %s.%s")
```


Here's the HTML needed in the R Markdown document to embed the GIF into the HTML file created with knitr.

```
<div id="PCA">  
    
</div>
```

Files: **Forensic-Glass-PCA.Rmd** and **Forensic-Glass-PCA.html**

Exploratory Analysis

Principal Component Analysis



WinF WinNF Veh Con Tabl Head

Files: **Forensic-Glass-PCA.Rmd** and **Forensic-Glass-PCA.html**

Take Home

- Variety of ways to select features for machine learning models.
- Explore several methods.
- Boruta often easy to apply without much data preparation. Many methods require pre-processing, like centering and scaling.
- Experiment only with training data to avoid “data leaks” and overfitting.